

systemd

Lucas Nussbaum

lucas.nussbaum@univ-lorraine.fr

Licence professionnelle ASRALL

Administration de systèmes, réseaux et applications à base de logiciels libres



UNIVERSITÉ
DE LORRAINE



nancy Charlemagne
Département Informatique

Outline

- 1 Introduction
- 2 Behind the scenes: cgroups
- 3 Managing services
- 4 Analyzing startup performance
- 5 Exploring the system status
- 6 Configuring services by writing unit files
- 7 Timer units
- 8 Socket activation
- 9 Logging with journald
- 10 Conclusions

Init system

- ▶ First process started by the kernel (pid 1)
- ▶ Responsible for **bringing up the rest of userspace**
 - ◆ Mounting filesystems
 - ◆ Starting services
 - ◆ ...
- ▶ Also the parent for orphan processes
- ▶ Traditional init system on Linux: **sysVinit**
 - ◆ Inherited from Unix System V
 - ◆ With additional tools (insserv, startpar) to handle dependencies and parallel initialization

systemd

- ▶ Written (since 2010) by Lennart Poettering (Red Hat) and others
- ▶ Now the default on most Linux distributions
- ▶ Shifts the scope from *starting all services* (sysVinit) to *managing the system and all services*
- ▶ Key features:
 - ◆ Relies on cgroups for
 - ★ Services supervision
 - ★ Control of services execution environment
 - ◆ Declarative syntax for unit files \leadsto more efficient/robust
 - ◆ Socket activation (on-demand startup of services)
 - ◆ Nicer user interface (systemctl & friends)
- ▶ Additional features: logging, timer units (cron-like), user sessions handling, containers management

Behind the scenes: cgroups

- ▶ Abbreviated from *control groups*
- ▶ Linux kernel feature
- ▶ Limit, account for and isolate **processes and their resource usage** (CPU, memory, disk I/O, network, etc.)
- ▶ Related to **namespace isolation**:
 - ◆ Isolate processes from the rest of the system
 - ◆ *Chroots on steroids*
 - ◆ PID, network, UTS, mount, user, etc.
- ▶ LXC, Docker \approx cgroups + namespaces (+ management tools)

cgroups and systemd

- ▶ Each service runs in its own cgroup
- ▶ Enables:
 - ◆ Tracking and killing all processes created by each service
 - ◆ Per-service accounting and resources allocation/limitation
- ▶ Previously, with sysVinit:
 - ◆ No tracking of which service started which processes
 - ★ PID files, or hacks in init scripts: `pidof` / `killall` / `pgrep`
 - ★ Hard to completely terminate a service (left-over CGI scripts when killing Apache)
 - ◆ No resources limitation (or using `setrlimit` (= `ulimit`), which is per-process, not per-service)
- ▶ Also isolate user sessions \leadsto kill all user processes (not by default)
- ▶ More information: Control Groups vs. Control Groups and Which Service Owns Which Processes?

systemd-cgls: visualizing the cgroups hierarchy

```
└─1 /sbin/init
  └─system.slice
    └─apache2.service
      └─1242 /usr/sbin/apache2 -k start
      └─9880 /usr/sbin/apache2 -k start
      └─9881 /usr/sbin/apache2 -k start
    └─system-getty.slice
      └─getty@tty1.service
        └─1190 /sbin/agetty --noclear tty1 linux
      └─getty@tty2.service
        └─24696 /sbin/agetty --noclear tty2 linux
    └─system-postgresql.slice
      └─postgresql@9.4-main.service
        └─1218 /usr/lib/postgresql/9.4/bin/postgres -D /var/lib/postgresql/9.4/main -c
        └─1356 postgres: checkpoint process
        └─1357 postgres: writer process
        └─1358 postgres: wal writer process
        └─1359 postgres: autovacuum launcher process
        └─1360 postgres: stats collector process
    └─gdm.service
      └─1209 /usr/sbin/gdm3
      └─1238 /usr/bin/Xorg :0 -novtswitch -background none -noreset -verbose 3 -auth ,
  └─user.slice
    └─user-1000.slice
      └─session-1.scope
        └─1908 gdm-session-worker [pam/gdm-password]
        └─1917 /usr/bin/gnome-keyring-daemon --daemonize --login
        └─1920 gnome-session
        └─1966 /usr/bin/dbus-launch --exit-with-session gnome-session
```

systemd-cgtop: per-service resources usage

Path	Tasks	%CPU	Memory	Input/s	Output/s
/	92	68.8	-	0B	243.9K
/system.slice	-	65.8	-	-	-
/system.slice/ModemManager.service	1	-	-	-	-
/system.slice/NetworkManager.service	2	-	-	-	-
/system.slice/accounts-daemon.service	1	-	-	-	-
/system.slice/apache2.service	3	0.1	-	-	-
/system.slice/atd.service	1	-	-	-	-
/system.slice/avahi-daemon.service	2	0.0	-	-	-
/system.slice/colord.service	1	-	-	-	-
/system.slice/system-postgresql.slice	8	66.0	-	340.4K	112.4M
/system.slice/system-postgresql.slice/postgresql@9.4-main.service	8	-	-	-	-
/system.slice/systemd-journald.service	1	-	-	-	-
/system.slice/systemd-logind.service	1	0.0	-	-	-
/user.slice	13	1.6	-	-	-
/user.slice/user-1001.slice	-	1.6	-	-	-
/user.slice/user-1001.slice/session-2.scope	4	-	-	-	-
/user.slice/user-1001.slice/session-4.scope	6	1.6	-	-	-
/user.slice/user-1001.slice/session-6.scope	5	-	-	-	-

Requires enabling CPUAccounting, BlockIOAccounting, MemoryAccounting

Managing services with systemctl

- ▶ What is being manipulated is called a *unit*: services (.service), mount points (.mount), devices (.device), sockets (.socket), etc.
- ▶ Basic commands:

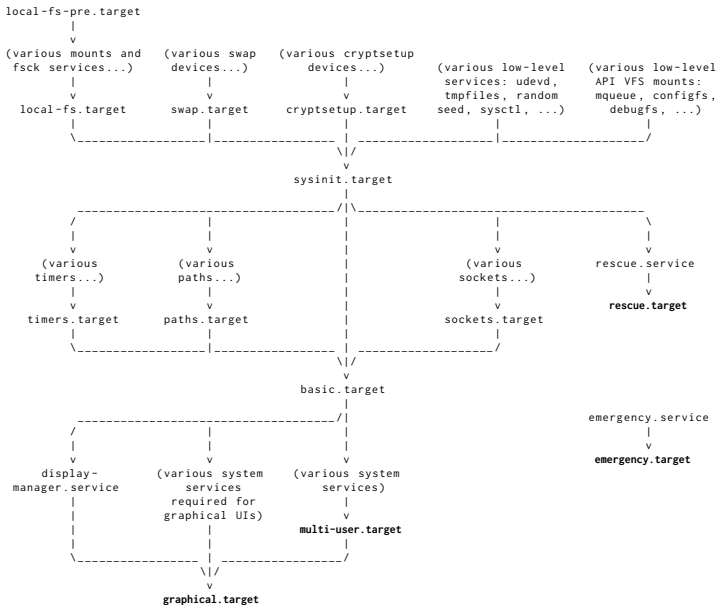
sysVinit	systemd	notes
service foo start service foo stop service foo restart service foo reload service foo condrestart update-rc.d foo enable update-rc.d foo disable	systemctl start foo systemctl stop foo systemctl restart foo systemctl reload foo systemctl condrestart foo systemctl enable foo systemctl disable foo systemctl is-enabled foo	restart if already running auto-start at next boot disable auto-start

- ▶ There's auto-completion (apache2 and apache2.service work)
- ▶ Several services can be specified:
systemctl restart apache2 postgresql

systemd and runlevels

- ▶ With sysVinit, runlevels control which services are started automatically
 - ◆ 0 = halt; 1 = single-user / minimal mode; 6 = reboot
 - ◆ Debian: no difference by default between levels 2, 3, 4, 5
 - ◆ RHEL: 3 = multi-user text, 5 = multi-user graphical
- ▶ systemd replaces runlevels with **targets**:
 - ◆ Configured using symlinks in `/etc/systemd/system/target.wants/`
 - ◆ `systemctl enable/disable` manipulate those symlinks
 - ◆ `systemctl mask` disables the service and prevents it from being started manually
 - ◆ The default target can be configured with `systemctl get-default/set-default`
 - ◆ More information: The Three Levels of "Off"

Default targets (bootup(7))



Analyzing startup performance

- ▶ Fast boot matters in some use-cases:
 - ◆ Virtualization, Cloud:
 - ★ Almost no BIOS / hardware checks \leadsto only software startup
 - ★ Requirement for infrastructure elasticity
 - ◆ Embedded world
- ▶ systemd-analyze time: summary

Startup finished in 4.883s (kernel) + 5.229s (userspace) = 10.112s

- ▶ systemd-analyze blame: worst offenders

```
2.417s systemd-udev-settle.service
2.386s postgresql@9.4-main.service
1.507s apache2.service
240ms NetworkManager.service
236ms ModemManager.service
194ms accounts-daemon.service
```


systemd-analyze critical-chain

- Shows services in the critical path

```
graphical.target @5.226s
├─multi-user.target @5.226s
│   └─exim4.service @5.144s +81ms
│       └─postgresql.service @5.142s +1ms
│           └─postgresql@9.4-main.service @2.755s +2.386s
│               └─basic.target @2.743s
│                   └─timers.target @2.743s
│                       └─systemd-tmpfiles-clean.timer @2.743s
│                           └─sysinit.target @2.742s
│                               └─networking.service @2.589s +153ms
│                                   └─local-fs.target @2.587s
│                                       └─run-user-117.mount @3.877s
│                                           └─local-fs-pre.target @223ms
│                                               └─systemd-remount-fs.service @218ms +4ms
│                                                   └─keyboard-setup.service @157ms +61ms
│                                                       └─systemd-udev.service @154ms +2ms
│                                                           └─systemd-tmpfiles-setup-dev.service @113ms +33ms
│                                                               └─kmod-static-nodes.service @102ms +10ms
│                                                                   └─system.slice @96ms
│                                                                       └─-.slice @94ms
```

Exploring the system status

- ▶ Listing units with `systemctl list-units` (or just `systemctl`):
 - ◆ active units: `systemctl`
 - ◆ List only services: `systemctl -t service`
 - ◆ List units in failed state: `systemctl --failed`
- ▶ Whole system overview: `systemctl status`
- ▶ GUI available: `systemadm`

systemctl status service

● avahi-daemon.service - Avahi mDNS/DNS-SD Stack

Loaded: loaded (/lib/systemd/system/avahi-daemon.service; enabled)

Active: active (running) since Wed 2015-04-01 21:49:28 CEST; 27s ago

Main PID: 2858 (avahi-daemon)

Status: "avahi-daemon 0.6.31 starting up."

CGroup: /system.slice/avahi-daemon.service

└─2858 avahi-daemon: running [grep.local]

└─2859 avahi-daemon: chroot helper

Apr 01 21:49:28 grep avahi-daemon[2858]: No service file found in /etc/avahi/services.

Apr 01 21:49:28 grep avahi-daemon[2858]: Joining mDNS multicast group on interface eth0.IPv

Apr 01 21:49:28 grep avahi-daemon[2858]: New relevant interface eth0.IPv6 for mDNS.

Apr 01 21:49:28 grep avahi-daemon[2858]: Joining mDNS multicast group on interface eth0.IPv

Apr 01 21:49:28 grep avahi-daemon[2858]: New relevant interface eth0.IPv4 for mDNS.

Apr 01 21:49:28 grep avahi-daemon[2858]: Network interface enumeration completed.

Apr 01 21:49:28 grep avahi-daemon[2858]: Registering new address record for fe80::d6be:d9ff

Apr 01 21:49:28 grep avahi-daemon[2858]: Registering new address record for 152.81.5.183 on

Apr 01 21:49:28 grep avahi-daemon[2858]: Registering HINFO record with values 'X86_64'/'LIN

Apr 01 21:49:29 grep avahi-daemon[2858]: Server startup complete. Host name is grep.local.

Includes:

- ▶ Service name and description, state, PID
- ▶ Free-form status line from `systemd-notify(1)` or `sd_notify(3)`
- ▶ Processes tree inside the cgroup
- ▶ Last lines from `journald` (syslog messages and stdout/stderr)

Configuring services by writing unit files

- ▶ With sysVinit: shell scripts in `/etc/init.d/`
 - ◆ Long and difficult to write
 - ◆ Redundant code between services
 - ◆ Slow (numerous `fork()` calls)
- ▶ With **systemd: declarative syntax** (.desktop-like)
 - ◆ Move intelligence from scripts to systemd
 - ◆ Covers most of the needs, but shell scripts can still be used
 - ◆ Can use includes and overrides (`systemd-delta`)
 - ◆ View config file for a unit: `systemctl cat atd.service`
 - ◆ Or just find the file under `/lib/systemd/system/` (distribution's defaults) or `/etc/systemd/system` (local overrides)

Simple example: atd

[Unit]

Description=Deferred execution scheduler

Pointer to documentation shown in systemctl status

Documentation=man:atd(8)

[Service]

Command to start the service

ExecStart=/usr/sbin/atd -f

IgnoreSIGPIPE=false # Default is true

[Install]

Where "systemctl enable" creates the symlink

WantedBy=multi-user.target

Common options

- ▶ Documented in `systemd.unit(5)` ([Unit]), `systemd.service(5)` ([Service]), `systemd.exec(5)` (execution environment)
- ▶ Show all options for a given service:
`systemctl show atd`
- ▶ Sourcing a configuration file:
`EnvironmentFile=-/etc/default/ssh`
`ExecStart=/usr/sbin/sshd -D $SSHD_OPTS`
- ▶ Using the `$MAINPID` magic variable:
`ExecReload=/bin/kill -HUP $MAINPID`
- ▶ Auto-restart a service when crashed: (\approx `runit` / `monit`)
`Restart=on-failure`
- ▶ Conditional start:
`ConditionPathExists=!/etc/ssh/sshd_not_to_be_run`
Conditions on architecture, virtualization, kernel cmdline, AC power, etc.

Options for isolation and security

- ▶ Use a **network namespace** to isolate the service from the network:
`PrivateNetwork=yes`
- ▶ Use a **filesystem namespaces**:
 - ◆ To provide a service-specific `/tmp` directory:
`PrivateTmp=yes`
 - ◆ To make some directories inaccessible or read-only:
`InaccessibleDirectories=/home`
`ReadOnlyDirectories=/var`
- ▶ Specify the list of **capabilities(7)** for a service:
`CapabilityBoundingSet=CAP_CHOWN CAP_KILL`
Or just remove one:
`CapabilityBoundingSet=~CAP_SYS_PTRACE`
- ▶ Disallow forking:
`LimitNPROC=1`

Options for isolation and security (2)

- ▶ Run as user/group: User=, Group=
- ▶ Run service inside a chroot:
RootDirectory=/srv/chroot/foobar
ExecStartPre=/usr/local/bin/setup-foobar-chroot.sh
ExecStart=/usr/bin/foobard
RootDirectoryStartOnly=yes
- ▶ Control CPU shares, memory limits, block I/O, swappiness:
CPUShares=1500
MemoryLimit=1G
BlockIOWeight=500
BlockIOReadBandwidth=/var/log 5M
ControlGroupAttribute=memory.swappiness 70
- ▶ More information: Converting sysV init scripts to systemd service files, Securing your services, Changing roots, Managing resources

Timer units

- ▶ Similar to cron, but with all the power of systemd (dependencies, execution environment configuration, etc)
- ▶ **Realtime (wallclock) timers**: calendar event expressions
 - ◆ Expressed using a complex format (see `systemd.time(7)`), matching timestamps like: `Fri 2012-11-23 11:12:13`
 - ◆ Examples of valid values: `hourly` (`= *-*-* *:00:00`), `daily` (`= *-*-* 00:00:00`), `*:2/3` (`= *-*-* *:02/3:00`)
- ▶ **Monotonic timers**, relative to different starting points:
 - ◆ 5 hours and 30 mins after system boot: `OnBootSec=5h 30m`
 - ◆ 50s after systemd startup: `OnStartupSec=50s`
 - ◆ 1 hour after the unit was last activated: `OnUnitActiveSec=1h` (can be combined with `OnBootSec` or `OnStartupSec` to ensure that a unit runs on a regular basis)

Timer units example

► myscript.service:

```
[Unit]
```

```
Description=MyScript
```

```
[Service]
```

```
Type=simple
```

```
ExecStart=/usr/local/bin/myscript
```

► myscript.timer:

```
[Unit]
```

```
Description=Runs myscript every hour
```

```
[Timer]
```

```
# Time to wait after booting before we run first time
```

```
OnBootSec=10min
```

```
# Time between running each consecutive time
```

```
OnUnitActiveSec=1h
```

```
Unit=myscript.service
```

```
[Install]
```

```
WantedBy=multi-user.target
```

Timer units example (2)

- ▶ Start timer:
`systemctl start myscript.timer`
- ▶ Enable timer to start at boot:
`systemctl enable myscript.timer`
- ▶ List all timers:
`systemctl list-timers`

Socket activation

- ▶ Similar to inetd:
 - ◆ The service is not started automatically at boot
 - ◆ systemd listens for connections on behalf of service
 - ◆ systemd starts the service to handle the connection
- ▶ Two use cases:
 - ◆ Services that are **seldomly used**: cups, sshd
 - ◆ During system boot, **parallelize startup** of services and delay synchronization until first interaction using sockets
- ▶ Same idea for dbus activation and fs activation (autofs-like)
- ▶ More information: Converting inetd Service, Socket Activation for developers (+ follow-up)

Socket activation example (sshd)

► sshd.socket:

[Unit]

Description=SSH Socket for Per-Connection Servers

[Socket]

ListenStream=22

Accept=yes

[Install]

WantedBy=sockets.target

► sshd@.service:

[Unit]

Description=SSH Per-Connection Server

[Service]

ExecStart=-/usr/sbin/sshd -i

StandardInput=socket

Socket activation example (sshd) (2)

- ▶ sshd@.service means that this is an **instanciated service**
- ▶ There's one instance of sshd@.service per connection:

```
# systemctl --full | grep ssh
sshd@172.31.0.52:22-172.31.0.4:47779.service    loaded active running
sshd@172.31.0.52:22-172.31.0.54:52985.service loaded active running
sshd.socket                                loaded active listening
```

- ▶ Instanciated services are also used by getty
 - ◆ See Serial console and Instanciated services

Logging with journald

- ▶ Component of systemd
- ▶ Captures syslog messages, kernel log messages, initrd and early boot messages, messages written to stdout/stderr by all services
 - ◆ Forwards everything to syslog
- ▶ Structured format (key/value fields), can contain arbitrary data
 - ◆ But viewable as syslog-like format with `journalctl`
- ▶ Indexed, binary logs
- ▶ Rotation handled transparently
- ▶ Can replace syslog (but can also work in parallel)
- ▶ Not persistent across reboots by default – to make it persistent:
`mkdir -p /var/log/journal`
- ▶ Can log to a remote host (with `systemd-journal-gateway`, not in Debian yet)

Example journal entry

```
_SERVICE=systemd-logind.service  
MESSAGE=User harald logged in  
MESSAGE_ID=422bc3d271414bc8bc9570f222f24a9  
_EXE=/lib/systemd/systemd-logind  
_COMM=systemd-logind  
_CMDLINE=/lib/systemd/systemd-logind  
_PID=4711  
_UID=0  
_GID=0  
_SYSTEMD_CGROUP=/system/systemd-logind.service  
_CGROUPS=cpu:/system/systemd-logind.service  
PRIORITY=6  
_BOOT_ID=422bc3d271414bc8bc95870f222f24a9  
_MACHINE_ID=c686f3b205dd48e0b43ceb6eda479721  
_HOSTNAME=waldi  
LOGIN_USER=500
```

Using journalctl

- ▶ View the full log: `journalctl`
- ▶ Since last boot: `journalctl -b`
- ▶ For a given time interval: `journalctl --since=yesterday`
or `journalctl --until="2013-03-15 13:10:30"`
- ▶ View it in the verbose (native) format: `journal -o verbose`
- ▶ Filter by systemd unit: `journalctl -u ssh`
- ▶ Filter by field from the verbose format:
`journalctl _SYSTEMD_UNIT=ssh.service`
`journalctl _PID=810`
- ▶ Line view (\approx `tail -f`): `journalctl -f`
- ▶ Last entries (\approx `tail`): `journalctl -n`
- ▶ Works with bash-completion
- ▶ See also: [Journald design document](#), [Using the Journal](#)

Containers integration

- ▶ General philosophy: **integrate management of services from machines (VMs and containers) with those of the host**
 - ◆ `systemd-machined`: tracks machines, provides an API to list, create, register, kill, terminate machines
 - ◆ `machinectl`: command-line utility to manipulate machines
 - ◆ other tools also have containers support:
 - ★ `systemctl -M mycontainer restart foo`
 - ★ `systemctl list-machines`: provides state of containers
 - ★ `journalctl -M mycontainer`
 - ★ `journalctl -m`: combined log of all containers
- ▶ `systemd` has its own mini container manager: `systemd-nspawn`
- ▶ Other virtualization solutions can also talk to machined
- ▶ More information: Container integration

More stuff

- ▶ New cross-distro configuration files: `/etc/hostname`,
`/etc/locale.conf`, `/etc/sysctl.d/*.conf`,
`/etc/tmpfiles.d/*.conf`
- ▶ Tools to manage hostname, locale, time and date: `hostnamectl`,
`localectl`, `timedatectl`
- ▶ Support for watchdogs
- ▶ Handling of user sessions
 - ◆ Each within its own cgroup
 - ◆ Multi-seat support
 - ◆ `loginctl` to manage sessions, users, seats
- ▶ systemd networking: `systemd-networkd`, `systemd-resolved`
 - ◆ Started as a way to improve networking management for containers

Conclusions

- ▶ systemd revisits the way we manage Linux systems
 - ◆ *If we redesigned services management from scratch, would it look like systemd?*
- ▶ For service developers: easier to support systemd than sysVinit
 - ◆ No need to fork, to drop privileges, to write a pid file
 - ◆ Just output logs to stdout (redirected to syslog, with priorities)
- ▶ Some parts still have rough edges, or are still moving targets, but are promising: journal, containers, networking
- ▶ systemd might not be the final answer, but at least it's an interesting data point to look at